

Baseline: 100% compute

History KV saved in CXL:
load + compute

